

THE ELK STACK

SANTIAGO JAVIER CALCAGNO

KARLSRUHE INSTITUTE OF TECHNOLOGY

SEMINAR „BIG DATA TOOLS“

WHAT IS THE ~~ELK~~ ELASTIC STACK?¹




LOGSTASH

Data processing



ELASTICSEARCH


Search engine



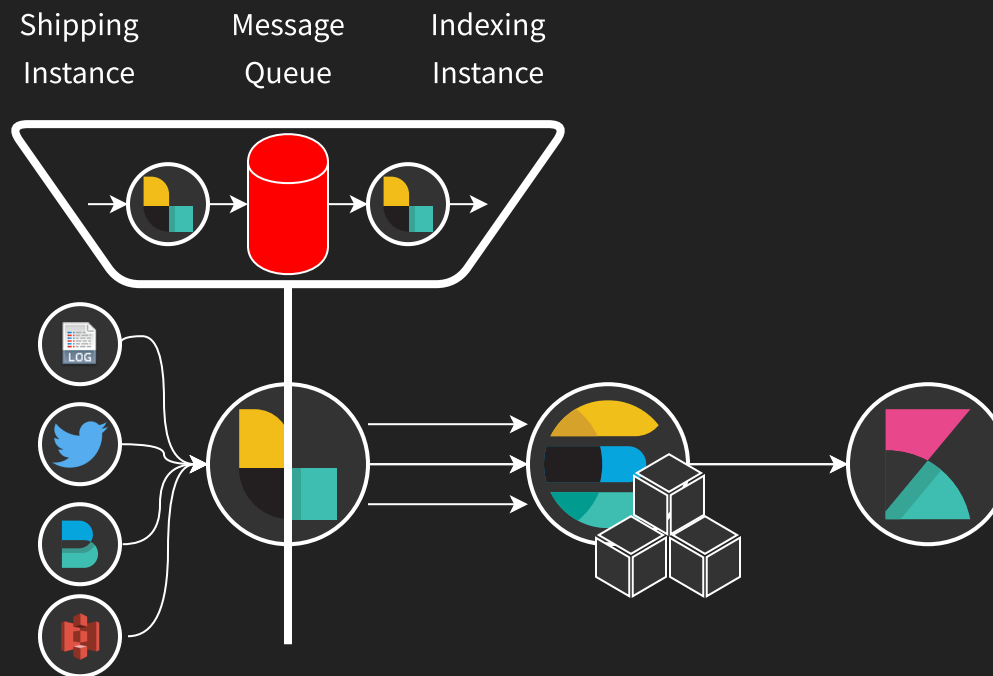
KIBANA

Visualization tool

GENERAL ARCHITECTURE



SCALING THE STACK²





Adobe

salesforce

XING



USE CASES

NETFLIX

More use cases [here.](#)³


ebay



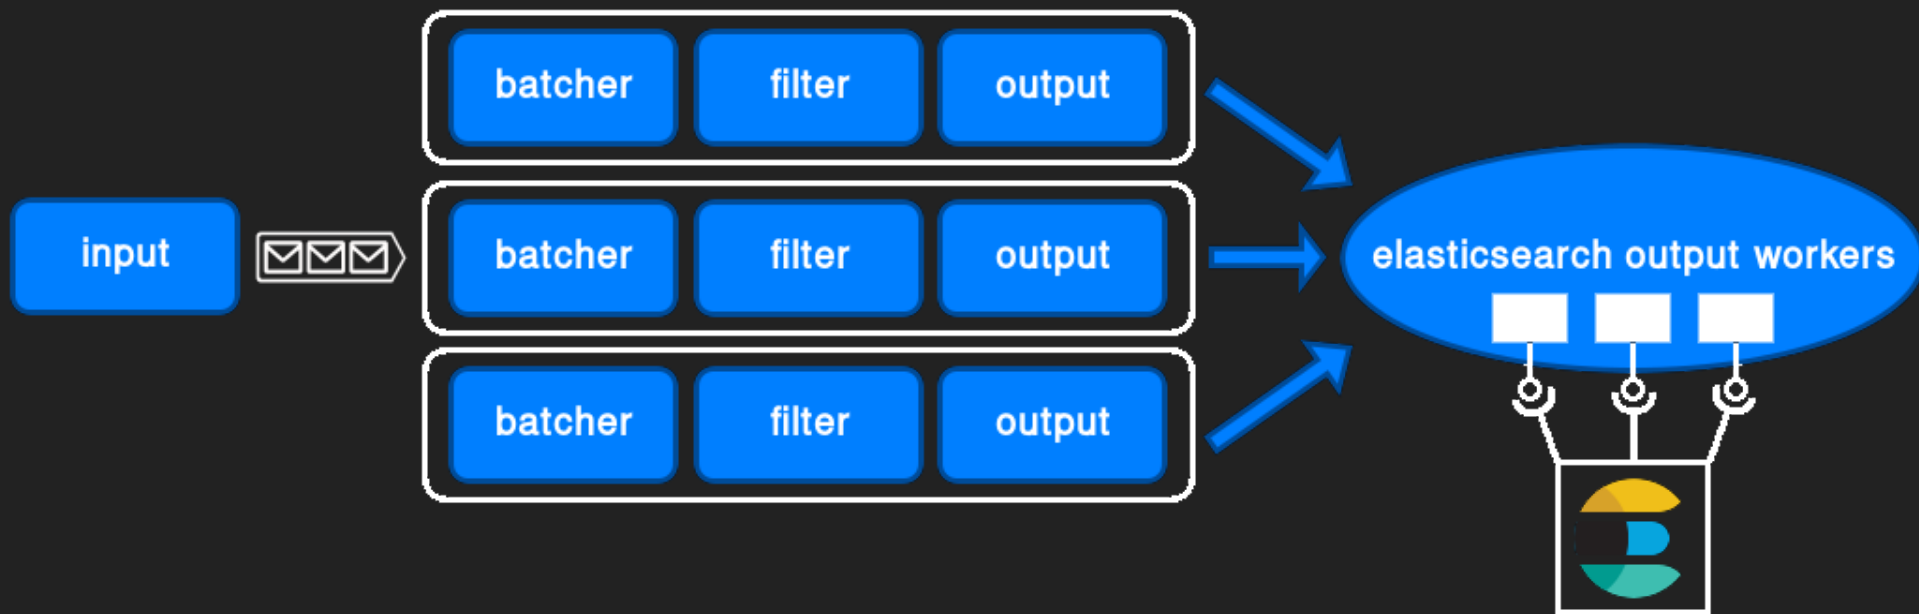
Goldman
Sachs




EVOLUTION OF THE LOGSTASH PIPELINE⁴



Logstash pipeline, versions 1.2.2 to 2.1



Logstash pipeline, version 2.2



Logstash pipeline, versions 5.0 and newer

GOAL

Analyze how the number of pipeline workers and the batch size affect the indexing rate (in a specific system).

- Intel® Core™ i5-2520M
- 16GB RAM DDR3-1866
- Samsung® EVO™ 250GiB mSATA SSD
- Arch Linux, kernel 4.8.13-1
- Elastic stack version 5.1

TESTS AND TOOLS


```
for w in "${W_VALUES[@]}"
do
    for b in "${B_VALUES[@]}"
    do
        sed -i -e "s/-w [0-9]*/-w $w/" docker-compose.yml
        sed -i -e "s/-b [0-9]*/-b $b/" docker-compose.yml
        docker-compose up &
        DOCKER_PID=$!
        sh ./gatherdata.sh &
        GATHER_PID=$!
        python jlog.py
        kill $GATHER_PID &&
        curl -s -XDELETE 'http://localhost:9200/_all'
        kill $DOCKER_PID
    done
done
```

run.sh

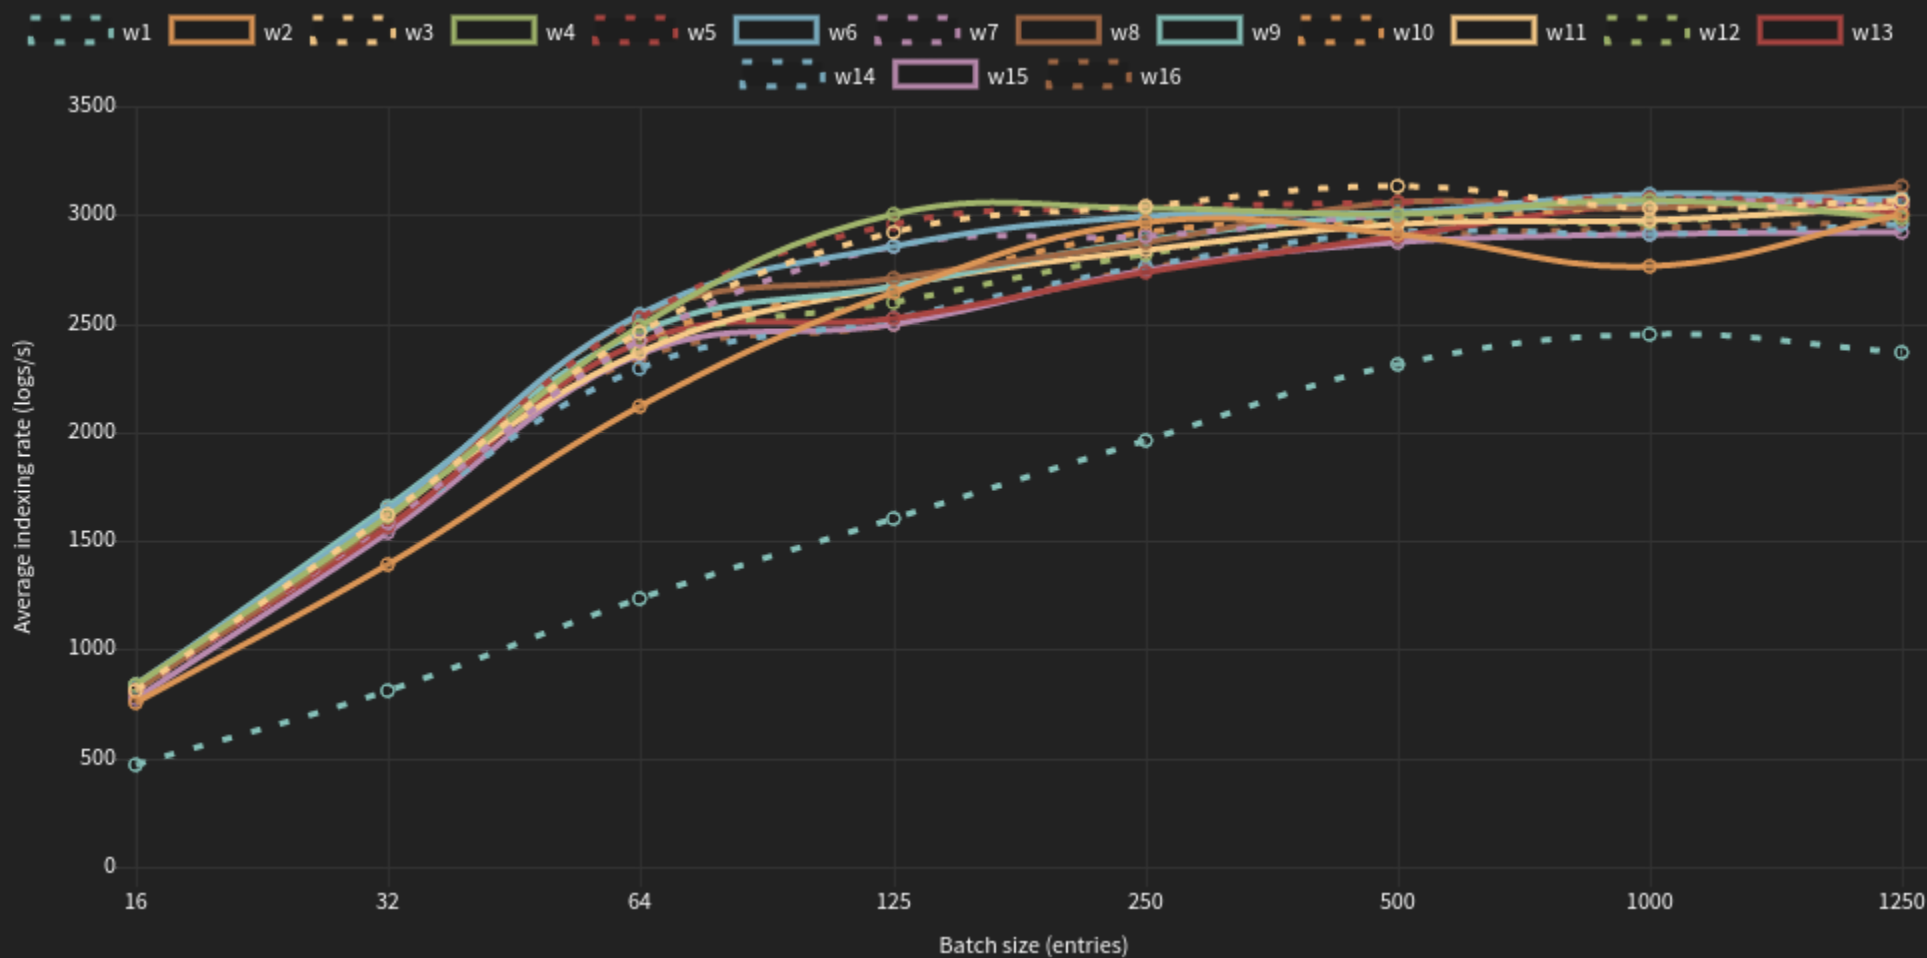
```
s.connect((TCP_IP, TCP_PORT))
for _ in range(0, MAX_LOGS):
    method = random.choice(methods)
    jlog = {
        'ip_src' : random.choice(ip_srcs),
        'websv' : random.choice(websvs),
        'method' : method,
        'query' : random.choice(gets) if method == 'GET'
                else random.choice(posts),
        'protocol' : random.choice(protocols),
        'response' : random.choice(responses),
        'user' : ''.join(random.choice(
            string.ascii_letters + string.digits)
            for _ in range(6)),
        'usertype' : random.choice(usertypes),
        'user_ip' : ".".join(map(str, (random.randint(0, 255)
            for _ in range(4)))),
    }
    msg = json.dumps(jlog) + '\n'
    s.send(msg.encode('utf-8'))
s.close()
```

jlog.py

RESULTS




Batch size comparison




Output workers comparison

SOME SPECIAL CASES



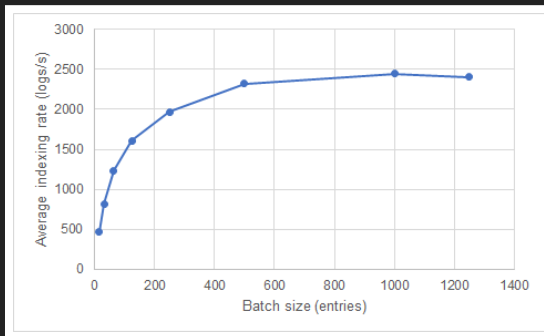
Batch size 16




Batch size 32




Batch size 125



1 worker



2 workers



4 workers

CONCLUSION

- For this system, $w \sim 4$ and $b \sim 150$
- Effect of message queue and more Elasticsearch nodes
- Generic testbed for more complex scenarios
- Try it yourself!

QUESTIONS?

THANK YOU!

References

1. Product Overview. (n.d.). Retrieved January 03, 2017, from <https://www.elastic.co/products>
2. Deploying and Scaling Logstash | Logstash Reference [5.1] | Elastic. (n.d.). Retrieved January 03, 2017, from <https://www.elastic.co/guide/en/logstash/current/deploying-and-scaling.html>
3. Use Cases. (n.d.). Retrieved January 08, 2017, from <https://www.elastic.co/use-cases>
4. Logstash Pipeline Architecture Discussion. (2016, July 21). Retrieved January 03, 2017, from <https://www.youtube.com/watch?v=FPLHS9Pmgk0>